

GlowProof

Judges' Quick-Test Guide

A 2-minute walkthrough of the live deployment for evaluators

Live URL: **<https://glow-proof.vercel.app>** — deployed in fixture mode, so nothing here spends real API credit or requires a webcam/real photo. Every score, mask, and product shown is real *code*, running against a real (cached) response shape — only the specific upstream call is swapped for a fixture.

The four things worth seeing (~2 minutes total)

1. Upload → measured scores (30s)

Click **"Upload a photo"** on the landing page and pick any image file.

Watch the staged progress ("Uploading over TLS" → "Locating facial landmarks" → ...) — this isn't decorative pacing, it mirrors the real Perfect Corp call's actual phases. Land on the results view: an overall score, a per-concern breakdown, and a colored region overlay that auto-plays once through every concern before settling on the worst one.

What to notice: click two or three concern cards directly — the overlay responds to your click, it's not just a fixed animation.

2. The routine never names a brand (30s)

Scroll to "Your routine." Each step names an action, a product *type* ("niacinamide serum"), and why — never a brand. That's enforced in the system prompt (`app/routine.py`), not just a formatting choice.

What to notice: the product card attached to each step (thumbnail, price, source store) is a *separate* system finding that product — the LLM never picked it. That split is the core of the whole pitch.

3. Recent scans — real persistence (30s)

Go back to the landing page. Below the claims list, a "Recent scans" panel lists past scans with a score and date. Click one — it reopens instantly from Xano, not from your browser's local state (open it in a fresh incognito window if you want to confirm that directly).

What to notice: this is a real database round trip, not a mock —

`api/history` and `api/scan/{id}` in the source both hit a live Xano API.

4. The disclaimer is permanent, not a footer (10s)

On the results view, note the cosmetic-guidance language throughout — this is enforced in the system prompt (`routine.py`'s `SYSTEM` constant forbids diagnosis language, disease names, and unhedged efficacy claims) and is visible UI copy, not a dismissible banner.

If you want to go deeper

- `/api/health` on the live URL shows exactly which integrations are live vs. fixture at that moment — no guessing about what's "real."
- `glowproof-technical-architecture.pdf` (submitted alongside this) covers the full system design, the Xano persistence architecture, the Vercel deployment, and a real bug the team caught during one deliberate live Perfect Corp capture — cited there as the argument for gating real calls behind a fixture flag rather than iterating against the live API.
- `VERIFIED.md` in the GitHub repo timestamps every claim above against a real, dated API response — nothing in this guide is asserted without a receipt somewhere in that file.